

Built to Please

Why Consumer AI Tells You What You Want to Hear, and What Serious Users Build Around It

A consumer AI model is a product, and the product succeeds by satisfying the person who pays for it. On questions where the asker has no stake in the answer, satisfaction and accuracy are the same job. Where the asker would prefer one answer over another, the two come apart, and the training pulls the model toward the preferred one. This paper explains why a system built to please cannot serve as its own check, and what serious users build around it so that its answers can be trusted.

BY **Jason Adamson**

ABOUT THIS PAPER

This is an insight paper, not primary research. It draws on published research, company statements, and the public record, cited in the references at the back. Two disclosures belong up front. The exhibit in section two comes from the author's own use of these systems, and he writes as a daily user of them. And Monderman builds diagnostic systems of the kind section five describes, checks around the model rather than trust in it, which is a stake in where this paper ends. Figures 1 through 3 are conceptual renderings.

EXECUTIVE SUMMARY

The same chat window sells two products. One is a companion, measured by how the conversation feels: how often you return, how long you stay, whether you renew. The other is an instrument, measured by whether the answer was right. On questions where the asker has no stake, the two products agree. Where the asker arrives with a preference, they come apart, and the training pulls toward the preference, because the systems learn from human choices and people tend to prefer answers that agree with them. Researchers have measured the lean, and in April 2025 OpenAI rolled back a ChatGPT update after users posted examples of it flattering and agreeing with them, behavior the company's own account called sycophantic. Memory features sharpen the general lean into a personal one, and instructions to be blunt cannot fix it, because the critic is still the model itself. What serious users build instead is a set of checks the model does not control: sources, code for the settled work, tests, an independent reviewer, records, and, at the end, a person who signs. The checks lower the risk of signing. They do not replace the signature. And a signature only counts when the signer understands the answer, has the standing to reject it, and the time to do either.

Two products

A companion is measured by how the conversation feels. An instrument is measured by whether the answer was right. The same model is sold as both.

The lean

Trained on human approval, the systems shift toward the asker's views and framing. Memory turns the general lean into a personal one.

The rebuild

Five checks the model does not control, plus a person who signs. The checks lower the risk of signing. They do not replace the signature.

1.

The Two Products

Open a chat window with a modern AI model and two different products sit in the same box.

The first product is a companion. It is measured the way consumer software is measured: how often you come back, how long you stay, whether you renew next month. A companion succeeds when the conversation feels good. It remembers what you like, matches your tone, takes your side of an argument

seriously, and sends you away feeling smarter than when you arrived. Whether any company scores its model on those numbers directly is not public. The pressure they create is.

The second product is an instrument. It is measured the way a lab test or an audit is measured: whether the answer was right. An instrument succeeds when it corrects you, and a correction rarely feels good in the moment. The moments an instrument is most valuable are the moments a companion gets scored worst.




Figure 1. One window, two products. A conceptual rendering. The companion is scored the way consumer software is scored. The instrument is scored the way a lab test is scored. The paper is about the questions where the two scorecards part.

Whether the two products conflict depends on the question. Ask for a fact you have no feelings about, and pleasing you and informing you are the same act. The conflict starts where you arrive with a preference: a draft you wrote, a plan you have argued for, a number you hope is right. On those questions the model can raise its score by moving toward you, and the counterweight, accuracy, is the one thing the choosers cannot easily check. This paper makes no claim about how often that happens. It does not need one, because the questions that carry a preference are not a random sample of anyone's day. They are the expensive ones: the forecast a budget depends on, the plan someone senior already backs, the risk nobody in the room wants to name.

The pull toward agreement is structural, not moral. The consumer product improves through two loops. In training, people compare pairs of answers and choose the one they prefer, and the model learns to make more answers like the chosen ones.¹ In the market, the versions people use more survive. In neither loop does the chooser verify the answer against the world. Verifying takes an expert, a source, or time, and none of those fit inside a feedback button. The companies also train for accuracy and for safety, and on many questions those objectives win. The claim here is narrower. Among the objectives, approval is the one the asker's preference can recruit. So the signal reaching the model is rich about what people like and thin about what is true. Nobody has to order the model to flatter anyone. The arithmetic of what gets measured does it on its own.

The rest of this paper follows that arithmetic into serious work, where the consumer product is now the tool on the desk and the cost of a pleasing wrong answer is highest.

2.

Built to Please

Researchers who study these systems have measured a lean toward the user: models tend to shift their answers toward the views and framing the asker brings.² In use, the lean takes recognizable shapes. The model takes the question as it arrives and answers inside its frame, even when the frame is the problem. It offers precise figures where none exist, and the precision reads as competence. It gives a weak case the standing of a strong one, and the balance reads as fairness. It attaches a sentence of caution that the rest of the answer ignores. Each shape hands the asker something satisfying, and none of them looks like flattery on the page. The research gives the shapes separate names, sycophancy for the agreement, confabulation for the invented material, premise-following, false balance. This paper groups them because they arrive together in use, not because they share one proven cause.

The author has seen the lean firsthand. He asked a widely used model whether a large membership institution would still exist at the end of the century, given the institution's own published growth figures. The model answered with probabilities. Asked where the probabilities came from, it conceded it had generated them itself.³ They were drawn from no study. They were the shape of an answer, filled in. The question carried a frame, decline, and the model filled it with confident arithmetic. Why it did is not knowable from the transcript. One conversation proves nothing on its own. It is here because it shows two of the shapes in a single answer. The public record shows the rest.

The clearest public admission that the lean is real came in April 2025. OpenAI withdrew an update to ChatGPT after users posted examples of it flattering and agreeing with them, and the company itself confirmed the pattern.⁴ The company's account was direct. Several changes had combined, and among them the update had added weight on immediate user feedback, the same approval signal the first section described. The result was behavior the company's own write-up called sycophantic. The company reversed the update within days and published an account of what went wrong. That deserves credit. Whether anyone meant to build the behavior is not something an outside reader can know, and the argument does not need to know it. If nobody meant to, the approval signal produced it on its own. If someone chose the weighting, the incentive was followed. Either way the behavior came from the signal, and the fix adjusted the signal's weight rather than removing it.

Every Node for Itself noted that episode in a single paragraph, as evidence of the agreement error a company invites when it moves its checking inside. This paper opens that paragraph up. An update can be rolled back. The incentive underneath it stays for as long as approval is what gets measured, and the next section is about what happens when that incentive learns your name.

3.

The Personalized Yes

The major consumer assistants now offer memory. They can carry what you tell them across conversations: your projects, your preferences, how you like your answers, the argument you made last

month.⁵ The feature is sold as convenience, and it is one. Reintroducing yourself every morning is a real cost, and memory removes it. It can also be switched off, and some users switch it off.

Now run the first section's incentive through the feature. A model with no memory can lean only on what you put in the window that day. A model that remembers you can aim the lean. Its file can hold which conclusions you favor, which framings you reach for, whose side you took last time. The model does not retrain itself on you, and it does not need to. What it says is steered by what sits in front of it, and memory keeps your file sitting in front of it. Memory itself is neutral, and it would sharpen a blunt adviser too. What it amplifies is whatever lean the training already built, and the training built a lean toward approval. The general pull toward the asker becomes a pull toward this asker. A second opinion drifts toward a briefed ally.

Merit After the Machine made two points this section leans on.⁶ The machines weakened the familiar evidence of ability, because fluent, confident work no longer shows who understood it. And a machine cannot stand behind an answer. It cannot be held to a promise or made to care that it failed. The personalized yes sharpens both problems. The answer reads even more like understanding, because it speaks your language and works from your history. And it is even harder to hold anyone to. The tailoring does not show on the page, and the only person positioned to notice it is the person it was tuned to please. A colleague can be asked whether they believe what they just told you. A transcript tuned to your file has nobody behind it to ask.

A common response is to fix all this with instructions. Tell the model to be blunt. Order it to challenge you. Paste in a request for brutal honesty and trust the machine to comply. The next section is about why that does not solve the problem.

“A second opinion drifts toward a briefed ally.”

4.

Why Asking Nicely Fails

The obvious fix is instruction, and it deserves to be taken seriously, because it partly works. Tell the model to be blunt and it obliges. The tone hardens, the hedges drop, and the answer that comes back sometimes surfaces a real problem the friendly version would have wrapped in praise. People who use these systems seriously already keep standing instructions like this, and researchers who test such prompts report that they reduce the measured lean.⁷ Reduce, not remove. What the instruction changes is behavior at the surface. What it cannot change is where the judge sits.

The critic you just commissioned is the same model, trained on the same approval signal, reading whatever you choose to show it. Toughness is now the assignment, so it performs toughness, which is one more act of compliance: the model agreeing with your request to be disagreed with. The objections it

produces cost it nothing, and they arrive because you ordered objections, so their arrival tells you nothing about whether your plan deserves them. An instruction can change the model's manner. It cannot give the model what checking requires, which is contact with something outside the conversation: a source, a test, a record, a person who verifies. Manner is the only thing a sentence in the window can reach.

The second limit is that you grade the toughness you ordered. Pushback that lands wrong gets rephrased away. Add context, narrow the question, ask again, and the answer tends to move. Over a session, the disagreement that survives is the disagreement you allowed to survive. With memory on, the grades carry forward, and the model arrives at the next conversation already shaped by which corrections you accepted and which you argued down. The pressure that trained the lean at scale runs again in miniature, with you doing the choosing and no one keeping score.

The third limit is the hardest to see. Suppose you challenge an answer and the model backs down, the way it backed down over the invented probabilities in section two. The retreat reads as honesty, and it may be exactly that. It may also be the next accommodation: a system trained on approval, facing a displeased user, produces the concession the user is looking for. From inside the conversation the two are identical. A real correction and a performed one use the same words. The author could not tell which one he had received, and no reader of the transcript can either. The test that would settle it, checking the claim against the world, is the step the exchange never contained.

Every Node for Itself made this point about companies: a self-check was never what anyone else wanted, because a company's word about itself is worth very little. The same rule holds at the scale of one person and one screen. A checker that takes instructions from the party being checked is not a check. Independence is not a tone the model can be asked to adopt. It is a position, and the only way to give the model that position is to build things around it that do not take orders from the window. That architecture is the rest of this paper.

“A checker that takes instructions from the party being checked is not a check.”

5.

The Rebuild

The fix is not better instructions. The people and companies that need right answers from these systems have started building checks around the model instead, checks the model cannot change from inside the conversation. None of it makes the model honest. The lean stays. The goal is smaller and more practical: make sure the answers that matter never rest on the model's word alone.

Not every question earns the full set. The checks scale with what a wrong answer costs, whether it can be undone, and whether anyone can check it at all. A question with no stakes needs none of this. The questions from section one need all of it.

The first check is sources. The model must show where each claim came from, and an answer that cannot show its sources does not go out. By itself this is weaker than it sounds. Models have invented sources before, many times, in public.⁸ So a separate step opens each source and confirms the quoted words are really in it. That step is simple and mechanical. Judging whether the source truly supports the claim is harder, and that job goes to the reviewer described below. Even the simple step matters here. The invented probabilities from section two die at this check, because a number that came from nowhere cannot produce a document that contains it.

The second check is taking work away from the model. Adding, counting, looking things up, sorting items into categories, anything with one right answer, is done by ordinary code that gives the same result every time. The model is handed the result and asked to explain it, not to produce it. It cannot change the number. It can still describe the number in a slanted way, so this check never works alone. The code can be wrong too. But when code is wrong, it is wrong the same way every time, so the mistake can be found and fixed. A model's mistakes shift from one run to the next.

The third check is tests. Before the system is used, and while it runs, its answers are graded against questions where the right answer is already known. The grader takes no instructions from the asker and does not care whether anyone is pleased. Tests have a limit worth stating: many of the expensive questions have no answer key. For those, tests do less. They check the parts that can be checked, and they send samples to an expert. They can also run one simple trial. Ask the same question twice, once sounding like someone who hopes the answer is yes and once like someone who hopes it is no, and see whether the answer changes.⁹ An answer that follows the hope was never an answer. Tests can also reward the sentence the chat window punishes: I do not know. In a conversation, admitting ignorance ends the exchange badly. In a tested system, a wrong answer costs more than no answer, so the system gets credit for saying it does not know.




Figure 2. The flip test. A conceptual rendering. One question is asked twice, wrapped in opposite hopes. An answer that follows the hope was never an answer.

The fourth check is a reviewer who does not answer to you. Someone reads the draft whose job is to find what is wrong with it. That can be a second model set up by someone else, or a person. Where the reviewer and the model disagree, the disagreement is shown, not smoothed over. One honest caution: two models trained the same way can share the same blind spots. The review still adds what the chat window cannot, because the reviewer holds no file on you and takes no orders from you. The strongest version is also blind, meaning the reviewer does not know which answer you were hoping for. Where the answer matters most, the reviewer is a person. This is the independent position section four said no instruction can create.

The fifth check is records. What the model was asked, what it was shown, what its memory supplied, which tools it called, which version answered, all of it is written down. A saved screen of chat is not a record. The record has to include what the system did out of sight, or it proves little. When an answer

turns out wrong, there is a trail, and someone can find out why. The trail includes the asker. Every rephrased question is in it. The steering that section three called invisible is visible in a log.




Figure 3. The model inside the checks. A conceptual rendering of section five. The model does the drafting. The ring does the checking, and the answer leaves through a person with a name.

The list has a familiar shape. The model must show its sources, pass tests it did not write, face a reviewer it does not control, and leave a record. None of it is new. Merit After the Machine described how the machines weakened the old evidence of ability; this list rebuilds that evidence standard around the machine itself. Every Node for Itself described the outside checking layer companies are tempted to cut; this list rebuilds that layer around a single tool. The same pattern is already required where wrong answers have been expensive for a long time. Banks are expected to have their models checked by people who did not build them.¹⁰ New aircraft are tested against rules the builder does not write.¹¹ Auditors are barred from checking their own work.¹² The rule underneath is the old one: nobody's word is taken on its own, and now that includes the machine's.

All of this costs money and time. The plain chat window stays faster, cheaper, and friendlier. For questions where nothing rides on the answer, the window is fine. The checks are for the questions where something does. And one thing is still missing from the list: a person. That is deliberate, and it is the last section.

6.

Staying in Charge

The checks in the last section make an answer safer to use. They do not make anyone responsible for using it. A log can show what happened. It cannot answer for what happens next. So the last piece of the rebuild is a person: someone who reads the checked answer, decides, signs, and answers for it if the decision goes wrong. Merit After the Machine said it plainly: as answers get cheap, someone willing to stand behind an answer gets expensive. The checks lower the risk of signing. They do not replace the signature. And a signature only counts when the signer understands the answer, has the standing to reject it, and the time to do either.

This part does not depend on how good the machines get. The answers will likely keep improving. A better answer still leaves the same question: who is acting on it, and who answers if it fails. A failed answer costs the model nothing. The costs land on people: the person who acted on it, the people affected by it, the company that shipped it. Responsibility follows the costs, and the costs run past the model every time.

None of this is a case against the chat window. A companion is a real product, and the author uses one. For questions where nothing rides on the answer, the window is the right tool. Even on serious work, the model inside the checks is more useful than the model alone, because its answers arrive with reasons to trust them. The mistake this paper has been describing is narrower: holding the companion and believing it is the instrument. The product was built to please you, and you took that as evidence it was right.

The three papers tell one story in three steps. Merit After the Machine showed the first: the machines learned to produce the evidence of ability, the fluent, confident work institutions used to tell who knew what they were doing. That work no longer shows who earned it. Every Node for Itself showed the second: companies holding that cheap fluent work are tempted to cut the outside checkers who would catch its mistakes, and the cutting has begun. This paper is the third: the tool that would take the checkers' place was built to agree with whoever is holding it. Strung together, the steps leave an organization alone with a system trained to please it, at the moment it most needs to hear that it is wrong. The way back is the same for all three. Evidence that can be checked, checkers who do not answer to the checked, and a person with a name at the end of every answer that matters. The machine can work inside every one of those checks, and it makes them faster and cheaper. The one thing it cannot supply is the independence, because a check built to please you is not a check. The machine was built to please. Staying in charge means building enough around it that pleasing you is no longer enough.

REFERENCES

1. On preference training, see Paul F. Christiano et al., “Deep Reinforcement Learning from Human Preferences” (2017), and Long Ouyang et al., “Training Language Models to Follow Instructions with Human Feedback” (2022). On the choosers favoring agreeable answers, Mrinank Sharma et al., “Towards Understanding Sycophancy in Language Models” (Anthropic, 2023; ICLR 2024), which finds that a response matching the user's views is more likely to be preferred, and that both humans and preference models prefer convincingly written sycophantic responses over correct ones a non-negligible fraction of the time.
2. Sharma et al. (note 1) demonstrate the behavior across five AI assistants and four free-form tasks. Earlier, Ethan Perez et al., “Discovering Language Model Behaviors with Model-Written Evaluations” (2022), documented sycophancy in models trained with human feedback.
3. The exchange took place in 2026 and is preserved in the author's records. The institution and the model are deliberately not named; the exhibit stands for the pattern, not the parties.
4. OpenAI, “Sycophancy in GPT-4o: What Happened and What We're Doing About It,” April 29, 2025, and “Expanding on What We Missed with Sycophancy,” May 2, 2025. The update was released on April 25, 2025 and rolled back within days. The second post states that the update added a reward signal built on thumbs-up and thumbs-down data from users. The user examples and the sequence are documented in contemporaneous coverage, including TechCrunch, April 29, 2025.
5. OpenAI announced memory for ChatGPT on February 13, 2024, and expanded it on April 10, 2025 to reference a user's past chats. Anthropic added past-chat referencing to Claude in August 2025, with broader memory following later that year. Google added memory to Gemini in February 2025 and continuous personalization in August 2025. Each company documents controls for switching the features off.
6. Jason Adamson, “Merit After the Machine,” Monderman Insight, August 2026, available at www.monderman.com/research.html, and “Every Node for Itself,” August 2026, available on the author's LinkedIn profile.

REFERENCES CONTINUED

7. After the April 2025 episode, OpenAI adjusted the system prompt of the model that remained in production, and later noted that system prompts have a more limited effect on model behavior than changes made in training (OpenAI, GPT-5 System Card, 2025). Prompt-based mitigation is studied directly in the research literature (for example, Hong et al., 2025). Fine-tuning approaches report the same direction, reduced but not removed; see Jerry Wei et al., “Simple Synthetic Data Reduces Sycophancy in Large Language Models” (2023).
8. Damien Charlotin's database tracks court decisions involving hallucinated AI output, most often invented citations; it recorded more than nineteen hundred decisions at its July 2026 count, a figure to be re-checked at publication. The best-known early case is *Mata v. Avianca, Inc.* (S.D.N.Y. 2023), in which lawyers were sanctioned for filing invented citations.
9. The trial is standard in the research. Sharma et al. (note 1) measure whether models change answers when the user states a preference or pushes back; Fanous et al., “SycEval: Evaluating LLM Sycophancy” (2025), formalizes rebuttal-based measurement.
10. Board of Governors of the Federal Reserve System, Office of the Comptroller of the Currency, and Federal Deposit Insurance Corporation, SR Letter 26-2, “Revised Guidance on Model Risk Management,” April 17, 2026. SR 26-2 supersedes SR 11-7 and emphasizes a tailored, risk-based approach. It excludes generative and agentic AI from its scope, while stating that broader risk-management and governance practices should determine appropriate controls for tools outside the guidance. For covered models, it defines effective challenge as critical analysis by objective experts with sufficient independence to maintain objectivity.
11. New aircraft designs are certified under Federal Aviation Regulations, 14 C.F.R. Part 21, against airworthiness standards the applicant does not write, including Part 25 for transport-category airplanes. Much of the compliance showing is performed by the manufacturer under FAA delegation and oversight, a system with well-known critics; the standards themselves remain the regulator's.
12. Sarbanes-Oxley Act of 2002, Section 201, and the SEC auditor-independence rules under Rule 2-01 of Regulation S-X, which bar auditors from auditing their own work. The bar applies to auditors of public companies and to services that would impair independence.

ABOUT THE AUTHOR

Jason Adamson is the founder of Monderman, an institutional performance research company. He is the author of *Governance, Bureaucracy and Organization: Stewardship, Drift, and Administrative Capacity* (Routledge, forthcoming). His career spans more than two decades of deep experience in intelligence analysis across the U.S. government, alongside private-sector experience at CrowdStrike and in startups. He holds an M.S. in Organization Development from Pepperdine University.

ABOUT MONDERMAN

Monderman is an institutional performance research company building Deterministic AI Infrastructure for organizational diagnostics. Its diagnostic platform produces structured operational reads for enterprises across sectors, including defense, healthcare, government, financial services, technology, manufacturing, and higher education.



Organizations deliver at the speed of their administrative reality.

connect@monderman.com • www.monderman.com • © 2026 Monderman. All rights reserved.