GOVERNANCE AND PERFORMANCE

The Unmeasured Layer

Administrative Reality and the Risk Standard Reporting Misses

Organizations deliver at the speed of their administrative reality. Boards, executives, and acquirers measure a great deal about the institutions they steward. The layer where strategies actually succeed or fail is the one their standard reporting does not integrate.

BY Jason Adamson

August 2026

1. The Blind Spot

Consider what a board sees. Financial statements, current and trended. Market position and competitive dynamics. Leadership assessments. Engagement and culture scores. Compliance posture, audit findings, risk registers. Operational dashboards with throughput, quality, and cost. An acquirer in diligence sees the same picture at higher resolution, plus a quality-of-earnings report and a synergy model. By any reasonable standard, the modern institution is intensively measured.

Now consider the pattern none of that measurement predicts. A strategy is approved, capital is committed, the leadership team is capable, and delivery stalls anyway. The failure literature on this is old, large, and remarkably stable. Bain reported in 2024 that 88 percent of business transformations fail to achieve their original ambitions. A Boston Consulting Group study of more than 850 companies put the success rate of digital transformations near 35 percent. The Corporate Leadership Council found only 34 percent of major change initiatives succeed, and KPMG found one in three chief executives reporting that their transformations failed to realize the expected value. The famous claim that 70 percent of change efforts fail has a genuinely contested lineage, as Mark Hughes documented in 2011, and the exact rate moves with how failure is defined. But every serious estimate lands in the same territory: falling short is the norm, not the exception.

Acquisitions show the same shape. The PwC and Mergermarket study Creating Value Beyond the Deal found 53 percent of acquirers underperforming their industry peers on total shareholder return over the twenty-four months after closing, and decades of academic work reviewed in Harvard Business Review place acquisition underperformance at 70 to 90 percent depending on how failure is defined. And when the postmortems run, on transformations and deals alike, they repeatedly identify integration, decision rights, accountability, operating-model design, and execution among the causes, alongside the strategic and technical ones. The difference is coverage. The strategic and technical causes are routinely measured going in. The institutional ones are assessed far less consistently and rarely integrated into a single structural-risk view.

2. The Layer Nobody Integrates

Call what those postmortems keep finding administrative reality: the institution as it actually operates, as distinct from the institution as it is documented. It is observable in four conditions. Decision friction: how long it actually takes a decision to move from raised to resolved, how many hands it passes through, and where it stalls. Accountability structure: where responsibility actually sits, which is often a different document from the organization chart. Key-person dependence: which individuals the institution silently routes around its formal structure to reach, because they are the ones who reliably deliver. And capacity drift: the slow divergence between the operating model as designed and the operating model as accumulated, one accommodation at a time.

It is worth being precise about how this layer relates to its neighbors. It is not engagement, which measures sentiment, and it is not leadership quality, which measures individuals. Culture is the nearer neighbor, and the line is analytic rather than exclusive: serious culture frameworks reach well beyond how people feel, into shared norms, habitual practices, and how things actually get done, and those observations overlap with pieces of administrative reality. The difference is the construct and the purpose. Culture work characterizes the social system, usually in order to change behavior. Administrative reality treats decision paths, accountability, dependence, and drift as structural conditions of the institution, held as a governance and risk matter, at the institution level, over time. An organization can have engaged people, healthy values, and capable leaders, and still be structurally unable to move a cross-functional decision in under a quarter. When it stalls, nothing was wrong with the people. Something was wrong with the structure they were operating inside, and filing that under a people heading is one reason it goes unmanaged.

> ***"Organizations deliver at the speed of their administrative reality."***

---

3. Why the Layer Goes Unreported

Run the audit instrument by instrument. Financial reporting measures outcomes, and outcomes lag: the numbers tell you what the institution did last quarter, not what it is currently capable of doing. Engagement surveys measure sentiment, and sentiment is not structure: people can feel fine inside an institution that cannot execute, and miserable inside one that can. Compliance and audit measure conformance to defined controls, and an institution can be perfectly compliant and structurally incapable at the same time, because no control objective asks whether the operating model still fits its purpose. Operational dashboards measure the activity of the processes that were designed in; they do not ask the meta-question of whether the design still matches the work. None of the four reports structure in operation.

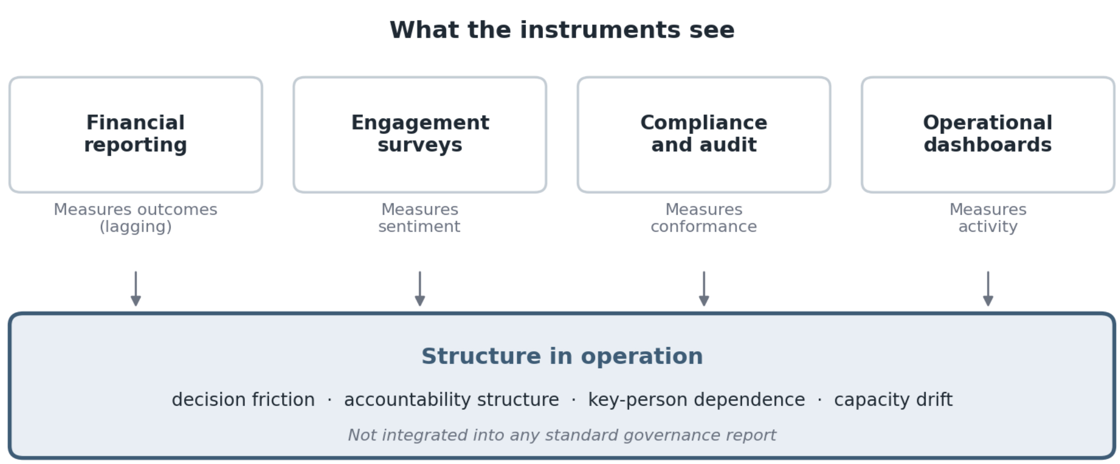

Figure 1. What the instruments see. Each standard instrument class answers the question it was built for. None of the four reports the structural layer where decision friction, accountability gaps, key-person dependence, and capacity drift accumulate, and standard governance reporting does not integrate that layer into a standing risk view.

To be fair to the field, specialized lenses exist that see fragments of this layer. Organizational network analysis, developed most prominently in Rob Cross's work, maps informal structure, hidden central actors, and collaboration bottlenecks, which is real visibility into dependence and friction. Process mining reconstructs how work actually moves through systems, deviations included. Internal audit tests whether controls operate as designed, and serious culture diagnostics observe behavioral norms. The blind spot, then, is not that nothing can see the layer. It is that these lenses are episodic, commissioned one at a time for particular questions, siloed by construct, and none of them feeds a standing, longitudinal, institution-level view of structural condition that reaches a board pack or a diligence file. Everything else in the governance stack was integrated into routine reporting long ago. This layer was assigned, by default, to intuition: leaders' feel for their institution, consultants' interviews, an occasional deep dive. All of that is real, and all of it shares the same three defects: it does not scale, it is not repeatable, and it leaves the building with the person who holds it.

4. How Drift Accumulates

The mechanism that makes this layer dangerous is that it decays through decisions that are each locally rational. A workaround is invented because the official process is slow. An extra approval is added after an incident, and never removed after the risk passes. A responsibility is informally absorbed by the person who happens to be good at it, and the formal owner quietly stops owning it. A standing meeting is created to compensate for a decision nobody is chartered to make. Every one of these accommodations is sensible on the day it is made. None of them triggers an alarm. The institution does not decide to drift; it accumulates drift out of reasonable choices.

The accumulation is not small. Gary Hamel and Michele Zanini estimated in Harvard Business Review that excess bureaucracy costs the U.S. economy more than $3 trillion a year in lost output, about 17 percent of GDP, and that nearly one in three American employees either works as a bureaucrat or spends the majority of their time on bureaucratic tasks. A Deloitte survey behind that estimate found non-administrative staff spending roughly 6.4 hours a week, about 16 percent of their time, complying with internal regulation. The decision-friction numbers are just as stark. Bain's research finds senior leaders and middle managers spending more than half their time in meetings, about two-thirds of meetings ending before the important decisions are made, and, in one company Bain studied, decisions being postponed roughly 60 percent of the time. Bain's Rule of Seven holds that every person added to a decision-making group beyond seven cuts decision effectiveness by about ten percent. In one firm, a single weekly leadership meeting consumed 300,000 hours a year across the organization once the preparation it demanded cascaded downward.

Key-person dependence grows by the same silent mechanism. Work routes toward whoever reliably delivers, until the organization chart and the load-bearing structure are two different documents, and the gap between them surfaces only on departure. The value at stake is measurable. In a PwC survey of 100 private equity executives, deals that lost 20 to 30 percent of key employees after close were associated with significant value destruction in 81 percent of cases, while deals that held key-employee turnover to 5 percent or below showed substantial value gains in 41 percent.

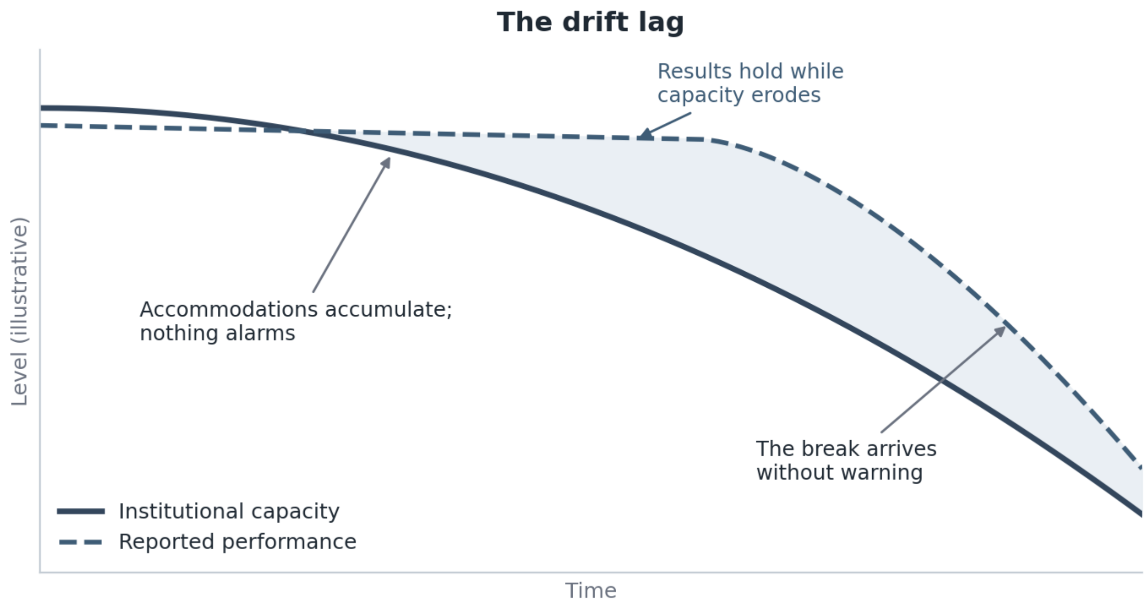

Figure 2. The drift lag. An illustrative mechanism sketch, not measured data. Reported performance runs on accumulated capacity, so results hold while the structure beneath them erodes. The gap between the curves is the risk class this paper describes: real, growing, and integrated into no standard report.

Figure 2 shows why the layer stays out of view until it does not. Results run on accumulated capacity, so reported performance holds while the structure beneath it erodes, and the instruments from Section 3 are pointed at the line that is still flat. When the break arrives, it arrives as a surprise, and the surprise gets attributed to whatever was nearest: a departure, a failed program, a bad quarter. The cause was upstream, years in the making, and unreported the entire time.

5. What Taking the Layer Seriously Would Require

If the layer is real and consequential, then treating it as a managed risk class implies requirements, stated here in principle rather than as any particular solution. First, observation of structure rather than sentiment: indicators anchored in how decisions, accountability, and dependence actually behave, not in how people feel about them. Second, longitudinal comparison: drift is a rate, and a rate is visible only against a baseline, which means snapshots, however sophisticated, cannot see it. Third, separation from individual-level constructs: the unit of analysis is the institution, not the employee, and instruments that grade people will keep landing in the HR file this layer must escape. Fourth, integration and validation: the fragments that specialized lenses already capture earn their place in a standing risk view only when they are assembled into one institution-level picture, and any claimed indicator of structural condition earns trust only by demonstrably predicting things that matter, such as decision cycle time, transformation milestones, unwanted departures, and overruns. Whether measurement meeting all four requirements exists yet in mature, validated form is a separate question from whether the layer exists. This paper argues the second claim, not the first.

The practical starting point costs nothing. A board, a chief executive, or an acquirer can begin by asking the structural questions the current information pack omits. How long does a cross-functional decision actually take here, and where did the last three stall? Which individuals, if they resigned tomorrow, would halt which processes, and does anything on paper acknowledge that? What has the operating model accumulated since it was designed, and when was any of it deliberately removed? The answers will be anecdotal at first. The point is that the questions belong on the standing agenda, in the governance file, where the layer they describe actually lives.

6. Where This Argument Could Be Wrong

Four objections deserve to be stated at full strength. First, existing constructs may capture more than this paper credits. Well-designed engagement and culture instruments correlate with performance across firms, and it is possible that sentiment and observed norms are a workable proxy for structure often enough to matter. The counter is that a cross-firm correlation is not visibility into a specific institution's structural condition, and the failure literature keeps citing structural causes in organizations whose people-scores were healthy. But the objection has real force.

Second, observable may not mean actionable. Even with the layer visible, accumulated structure defends itself; Hamel and Zanini's work documents how reliably bureaucracy regenerates after pruning. Measurement without a governance owner changes nothing, and it is fair to ask whether owners would act on structural findings any more readily than they act on the failure statistics already available. Third, judgment may substitute. Seasoned operators read institutions well, and for a single organization with a long-tenured, honest leadership team, cultivated judgment may outperform any instrument. The limits of judgment are scale, comparability, and portability: it cannot cover a portfolio, cannot be compared across units, and departs with its holder. For a single institution, though, the objection stands.

Fourth, the causal arrow could run backward. Perhaps performance trouble produces structural symptoms, rather than structural decay producing performance trouble, and the drift lag of Figure 2 has the sequence reversed. Settling that question requires longitudinal evidence linking measured structural conditions to subsequent outcomes, and that evidence base, in public and mature form, largely does not yet exist. Readers should hold the thesis with exactly that reservation.

7. Implications by Seat

For the board: structural condition belongs on the risk register, alongside cyber, liquidity, and concentration, because it behaves like those risks: it accumulates quietly, it is cheap to examine and expensive to discover, and it is the board's to steward because no operating owner naturally holds it. The structural questions in Section 5 are a sufficient annual starting point.

For the chief executive: the operating model is an asset that depreciates, and almost no organization schedules its maintenance. Capacity gets rebuilt in crisis, after the break in Figure 2, at the moment repair is most expensive. Treating structural review the way capital planning is treated, as a recurring discipline rather than a response, is the executive translation of this paper.

For the acquirer and the operating partner: diligence prices a great deal, but the institutional causes the postmortems keep naming are often not priced with the same discipline. The first hundred days run on the target's administrative reality, not on the deal model, and the base rates in Section 1 are the cost of learning that after close. The cheapest moment to see the layer is before signing: map how decisions actually move in the target, map who the institution actually depends on, and price what the map shows. For the transformation leader: a transformation is an attempt to change this layer using this layer, which is why the recurring causes in the failure literature are the layer itself. Sequencing structural repair ahead of capability building is not a delay to the program; it is the program.

Conclusion

The claim here is deliberately narrow. A layer of institutional conditions plausibly sits upstream of delivered performance. It is observable in decision friction, accountability structure, key-person dependence, and capacity drift. Fragments of it are visible to specialized disciplines, and it is analytically distinct from culture, engagement, and leadership. What standard governance reporting does not do is integrate those fragments into a standing, longitudinal, institution-level view, which leaves a consequential class of institutional risk carried largely blind. Institutions that learn to see their administrative reality will get the chance to steward it. The rest will keep reading dashboards that are working exactly as designed, measuring what they were built to measure, right up until the layer underneath them gives way.

About this paper: this is a synthesized argument drawing on published research, attributed in-line and listed below. Figure 2 is an illustrative mechanism sketch, not measured data. The failure-rate literature it cites is contested in its exact figures; the paper relies on the direction of that literature, not any single number.

Jason Adamson writes on governance, bureaucracy, and institutional performance. His book, Governance, Bureaucracy and Organization: Stewardship, Drift, and Administrative Capacity, is forthcoming from Routledge (www.routledge.com/9781041403296).

References

Bain & Company. Transformation research, 2024 (88 percent of business transformations fail to achieve original ambitions).

Bain & Company, Decision Insights series (Blenko, Mankins, Rogers): decision effectiveness, meeting practices, and the Rule of Seven.

Mankins, M., C. Brahm, and G. Caimi. Your Scarcest Resource. Harvard Business Review, 2014.

Boston Consulting Group. Digital transformation success research (study of more than 850 companies).

CEB Corporate Leadership Council. Change initiative success research, 2016.

Cross, R. What Is Organizational Network Analysis? robcross.org.

Deloitte. Get Out of Your Own Way: Unleashing Productivity. 2014.

Errida, A., and B. Lotfi. The Determinants of Organizational Change Management Success. International Journal of Engineering Business Management, 2021.

Hamel, G., and M. Zanini. Excess Management Is Costing the U.S. $3 Trillion Per Year. Harvard Business Review, September 2016; and The $3 Trillion Prize for Busting Bureaucracy, 2016.

Harvard Business Review. Don't Make This Common M&A Mistake. 2020.

Hughes, M. Do 70 Per Cent of All Organizational Change Initiatives Really Fail? Journal of Change Management, 2011.

KPMG. Global Transformation Study, 2016.

Knowledge at Wharton. Why Many M&A Deals Fail, and How to Beat the Odds. 2025.

PwC and Mergermarket. Creating Value Beyond the Deal, 2019 (acquirer TSR underperformance).

PwC. Creating Value Beyond the Deal: Private Equity, 2019 (survey of 100 private equity executives on key-employee turnover and deal value).

van der Aalst, W. Process Mining: Data Science in Action. Springer, 2016.